

```
In [1]: 1 # Import standard packages
2 import pandas as pd
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6 %matplotlib inline
7 sns.set_theme(style="darkgrid")
8 from collections import Counter
9 from sklearn.preprocessing import OneHotEncoder, LabelEncoder
10 import time
11 # from sklearn import metrics
12 from sklearn.metrics import precision_score, recall_score, accuracy_score
13 from sklearn.pipeline import Pipeline
14 from sklearn.preprocessing import MinMaxScaler, StandardScaler
15 from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV
16 from sklearn.metrics import confusion_matrix, plot_confusion_matrix
17
18 from sklearn.linear_model import LogisticRegression
19 from sklearn.ensemble import RandomForestClassifier
20 from xgboost import XGBClassifier
```

Tanzania Machine Learning Water Pump Classification

Modeling Notebook

Author: Dylan Dey

This project is available on github here:

https://github.com/ddey117/Tanzanian_Water_Pump_Classification
(https://github.com/ddey117/Tanzanian_Water_Pump_Classification)

The Author can be reached at the following email: ddey2985@gmail.com
(<mailto:ddey2985@gmail.com>)

Blog: *insert blog link*

Classification Metric Understanding

Below is a confusion matrix that would be produced from a model performing predictive maintenance on behalf of the Ministry of Water. There are four possible outcomes to be considered. The confusion matrix below is a visual aid to help in understanding what classification metrics to consider when building the model.

| ACTUAL | Predicted | |
|----------------|--|--|
| | Non-Functional | Functional |
| Non-Functional | True Negatives A true negative for my model would be a non-functional pump correctly labeled as a non-functional pump. Teams and/or additional resources will need to be sent here first. | False Positive A false positive in my model would be a non-functional pump incorrectly labeled as a functional pipe. False positives should be reduced as much as possible as this is the worst case scenario for my model. Teams/resources would be withheld from communities that need them. |
| Functional | False Negative A false negative for my model would happen when the model predicts a well that is actually functional to be non-functional. Therefore, teams/resources would be sent to communities that already have a functional water pump. Reducing false positives would help efficiency of distributing resources appropriately.. | True Positives A true positive for my model will be considered a functional pump correctly labeled as functional. This means no teams and/or additional resources will need to be sent here. |

A true positive in the current context would be when the model correctly identifies a functional pump as functional. A true negative would be when the model correctly identifies a non-functional pump as non-functional. Both are important and both can be described by the overall **accuracy** of the model.

True negatives are really at the heart of the model, as this is the situation in which the Ministry of Water would have a call to action. An appropriately outfitted team would be set to *all* pumps that my model identifies as non-functional. Thus, this is the situation in which the correct resources are being derived to the correct water pumps as quickly as possible. High accuracy would mean that more resources are going to the correct locations from the get-go.

True positives are also important. This is where the model will really be saving time, resources, and money for the Ministry of Water. *Any* pumps identified as functional would no longer need to be physically checked and the Ministry of Water can withhold additional resources from going to pumps that do not actually need them.

Notice the emphasis on *any* and *all* pumps in my description of true negatives and true positives above. The true cost/resource analysis is really the consideration of this fact: no model I create will ever correctly identify every single pump appropriately. This is the cost of predictive maintenance and a proper understanding of false positives and false negatives is extremely important in production of classification models in the given context.

False positives in the current context are the worst case scenario for modeling. This is the scenario in which the model **incorrectly** identifies a non-functional model as functional. Thus, resources would be withheld and no team would be sent to physically check these pumps, as the Ministry of Water would have to assume they are indeed functional if they want to use the model appropriately. False positives therefore describe the number of non-functional pumps that will go unvisited and unfixed until they can be resolved by other means. Reducing false positives as much as possible is very important.

Well, why would I want to build the model if these false positives cannot be completely avoided? Cost/resource management, of course! Afterall, it is about making sure as many people get clean water as quickly as possible. The reality is there that resources are finite, and without the model the Ministry of Water likely would not have the resources to physically check all the pumps and then fix all of the pumps in any sort of reasonable timeline, and even less communities would have access to fresh water when compared to using the model for predictive maintenance.

False negatives are also important to consider. While false positives can be considered more harmful overall, false negatives are also important to reduce as much as possible. In the given context, false negatives describe the situation in which the model **incorrectly** identifies a functional pump as non-functional. Because the Ministry of water will deploy fully equipped teams to visit all pumps that my model predicts to be non-functional, these will be the pumps that will waste resources. Resources will be sent to locations that they aren't needed, and the metric that describes this would be false negatives. Thus, reduction of false negatives is essential in improving the efficiency of resource management through predictive maintenance.

In summary, overall accuracy of the model and a reduction of both false negatives and false positives are the most important metrics to consider when developing a model in this context. More specifically, models will be tuned to **maximize accuracy and f1-score**.

Accuracy: The number of correct predictions made by model divided by the total number of predictions. This measures how successful my model is at labeling functional pumps as functional and non functional pumps as non functional. It is a very important metric to consider.

f1-score: I stated above that both false negatives and false positives are important to avoid. This means that I really want precision and recall to go be high and therefore false negatives and false positives to be low. The f1-score is the harmonic mean of precision and recall. Maximizing for this metric is the best way to get a balance of both false positives and false negatives.

Function Definition

Below are all of the functions used for preprocessing data before modeling.

```

In [2]: 1 def drop_cols(water_pump_df):
2         to_drop_final = ['id', 'recorded_by', 'num_private',
3                           'waterpoint_type_group', 'source',
4                           'source_class', 'extraction_type',
5                           'extraction_type_group', 'payment_type',
6                           'management_group', 'scheme_name',
7                           'water_quality', 'quantity_group',
8                           'scheme_management', 'longitude',
9                           'latitude', 'date_recorded',
10                          'amount_tsh', 'gps_height',
11                          'region_code', 'district_code']
12         #'population'
13
14         return water_pump_df.drop(columns=to_drop_final, axis=1)
15
16     #helper function to bin construction year
17 def construction_wrangler(row):
18     if row['construction_year'] >= 1960 and row['construction_year'] <
19         return '60s'
20     elif row['construction_year'] >= 1970 and row['construction_year']
21         return '70s'
22     elif row['construction_year'] >= 1980 and row['construction_year']
23         return '80s'
24     elif row['construction_year'] >= 1990 and row['construction_year']
25         return '90s'
26     elif row['construction_year'] >= 2000 and row['construction_year']
27         return '00s'
28     elif row['construction_year'] >= 2010:
29         return '10s'
30     else:
31         return 'unknown'
32
33 def bin_construction_year(water_pump_df):
34     water_pump_df['construction_year'] = water_pump_df.apply(lambda row:
35     return water_pump_df
36
37
38     #takes zero placeholders and NAN values and converts them into 'unknown'
39 def fill_unknowns(water_pump_df):
40     installer_index_0 = water_pump_df['installer'] == '0'
41     funder_index_0 = water_pump_df['funder'] == '0'
42     water_pump_df.loc[installer_index_0, 'installer'] = 'unknown'
43     water_pump_df.loc[funder_index_0, 'funder'] = 'unknown'
44     water_pump_df.fillna({'installer': 'unknown',
45                           'funder': 'unknown',
46                           'subvillage': 'unknown'}, inplace=True)
47     return water_pump_df
48
49     #returns back boolean features without NANs while maintaining same rat
50 def fill_col_normal_data(water_pump_df):
51     filt = water_pump_df['permit'].isna()
52     probs = water_pump_df['permit'].value_counts(normalize=True)
53     water_pump_df.loc[filt, 'permit'] = np.random.choice([True, False]
54                                                           size=int(filt.sum()),
55                                                           p = [probs[True], probs[False]])
56     filt = water_pump_df['public_meeting'].isna()

```

```

57     probs = water_pump_df['public_meeting'].value_counts(normalize=True)
58     water_pump_df.loc[filt, 'public_meeting'] = np.random.choice([True
59         size=int(filt.sum()),
60         p = [probs[True], probs[False]])
61     return water_pump_df
62
63
64
65 def apply_cardinality_reduct(water_pump_df, reduct_dict):
66     for col, categories_list in reduct_dict.items():
67         water_pump_df[col] = water_pump_df[col].apply(lambda x: x if x
68     return water_pump_df
69
70
71
72 #one_hot_incode categorical data
73 def one_hot(water_pump_df):
74     final_cat = ['funder', 'installer', 'wpt_name', 'basin', 'subvilla
75         'lga', 'ward', 'public_meeting', 'permit', 'construction_year',
76         'extraction_type_class', 'management', 'payment', 'quality_grou
77         'quantity', 'source_type', 'waterpoint_type']
78
79     water_pump_df = pd.get_dummies(water_pump_df[final_cat], drop_firs
80
81     return water_pump_df
82
83
84
85 #master function for cleaning dataframe
86 def clean_dataframe(water_pump_df, reduct_dict):
87     water_pump_df = drop_cols(water_pump_df)
88     water_pump_df = bin_construction_year(water_pump_df)
89     water_pump_df = fill_unknowns(water_pump_df)
90     water_pump_df = fill_col_normal_data(water_pump_df)
91     water_pump_df = apply_cardinality_reduct(water_pump_df, reduct_dic
92     water_pump_df = one_hot(water_pump_df)
93
94     return water_pump_df
95
96 #####
97 # The rest of the functions in this section
98 #define functions that reduce cardinality
99 #by mapping infrequent values ot other
100 #the dictionary derived from these functions
101 #will be used by my_funk in my master
102 #clean_dataframe function
103
104 #helper function for reducing cardinality
105 def cardinality_threshold(column, threshold=0.65):
106     #calculate the threshold value using
107     #the frequency of instances in column
108     threshold_value=int(threshold*len(column))
109     #initialize a new list for lower cardinality column
110     categories_list=[]
111     #initialize a variable to calculate sum of frequencies
112     s=0
113     #Create a dictionary (unique_category: frequency)

```

```

114     counts=Counter(column)
115
116     #Iterate through category names and corresponding frequencies after
117     #by descending order of frequency
118     for i,j in counts.most_common():
119         #Add the frequency to the total sum
120         s += dict(counts)[i]
121         #append the category name to the categories list
122         categories_list.append(i)
123         #Check if the global sum has reached the threshold value, if s
124         if s >= threshold_value:
125             break
126         #append the new 'Other' category to list
127         categories_list.append('Other')
128
129     #Take all instances not in categories below threshold
130     #that were kept and lump them into the
131     #new 'Other' category.
132     new_column = column.apply(lambda x: x if x in categories_list else
133 #         return new_column
134     return categories_list
135
136     #reduces the cardinality of appropriate categories
137 def get_col_val_mapping(water_pump_df):
138     col_threshold_list = [
139         ('funder', 0.65),
140         ('installer', 0.65),
141         ('wpt_name', 0.15),
142         ('subvillage', 0.07),
143         ('lga', 0.6),
144         ('ward', 0.05)
145     ]
146
147     reduct_dict = {}
148
149     for col, thresh in col_threshold_list:
150         reduct_dict[col] = cardinality_threshold(water_pump_df[col],
151                                                 threshold= thresh)
152
153     return reduct_dict
154
155     # reduct_dict is a key value mapper that will
156     # be used for both training and testing sets
157     # in order to reduce cardinality of the data

```

Import The Data From Multiple Sources

I used a number of sources for my data to use for modeling in this notebook. The cell below imports the original data from the DrivenData competition, data derived from DrivenData in QGIS and opensource hydrology data, and population data from Tanzania government census in 2012.

```

In [3]: 1 #import data from DrivenData
2 train_labels = pd.read_csv('data/0bf8bc6e-30d0-4c50-956a-603fc693d966.c
3 train_features = pd.read_csv('data/4910797b-ee55-40a7-8668-10efd5c1b960
4 df = train_features.merge(train_labels, on='id').copy()
5 #import QGIS derived data and prepare for model
6 river_df = pd.read_csv('data/river_dist2.csv')
7 #removing outliers
8 index_riv = river_df[river_df['HubDist'] > 66].index
9 river_median = river_df['HubDist'].median()
10 river_df.loc[index_riv, 'HubDist'] = river_median
11
12 #create boolean for pump being within 8 km of river
13 river_s = river_df['HubDist'].copy()
14 river_s.rename('near_river', inplace=True)
15 near_river = river_s[river_s < 8].apply(lambda x: 1 if not pd.isnull(x)
16 df = df.join(near_river)
17 df.near_river.fillna(0, inplace=True)
18
19 #import population data from 2012 government census
20 df_pop = pd.read_excel('data/tza-pop-popn-nbs-baselinedata-xlsx-1.xlsx')
21
22 #create a dictionary of values with format {Ward : Total Population}
23
24 pop_index = df_pop.groupby('Ward_Name')['total_both'].sum().index
25 pop_values = df_pop.groupby('Ward_Name')['total_both'].sum().values
26 pop_dict = dict(zip(pop_index, pop_values))
27
28 #create pandas Dataframe for merging
29 pop_dataframe = pd.DataFrame.from_dict(pop_dict, orient='index')
30 #rename column for clarity
31 pop_dataframe.rename(columns={0: 'ward_pop'}, inplace=True)
32
33 #merge dataframes
34 df_pop_merge = df.merge(pop_dataframe,
35                          how='left',
36                          left_on='ward',
37                          right_index=True)
38
39 #replace null values of ward population with
40 #median ward population
41
42 ward_pop_median = df_pop_merge['ward_pop'].median()
43 df_pop_merge.fillna(value=ward_pop_median, inplace=True)
44 # merge back into df and drop pop column
45 ward_pop_s = df_pop_merge['ward_pop'].copy()
46 df = df.join(ward_pop_s)
47 df.drop(columns=['population'], axis=1, inplace=True)
48
49 #bin functional needs repair and non function together
50 #simplify problem into binary classification
51 need_repair_index = df['status_group'] == 'functional needs repair'
52 df_binary = df.copy()
53 df_binary.loc[need_repair_index, 'status_group'] = 'non functional'
54
55 #workaround for Label Encoding mapping
56 #Label Encoding doesn't have a mapper method

```



```

57 #but it seems to encode alphabetically
58 #I want Non functional to be zero but encoded as 1
59 #unless make it come before 'functional' alphabetically
60 funct_index = df_binary['status_group'] == 'functional'
61 nonfunct_index = df_binary['status_group'] == 'non functional'
62 df_binary.loc[funct_index, 'status_group'] = 'Zfunctional'
63 df_binary.loc[nonfunct_index, 'status_group'] = 'Anon functional'

```

Transform Target to Numerical Data with Label Encoding

As shown at the beginning of the project, the functional state of the water supply point is described as: functional — It is working; non functional — it is not working; functional needs repair — Running, but needing maintenance.

These groups will be relabeled as: functional: 0 functional needs repair: 1 non functional: 2

This would be for future work in which I didn't bin functional needs repair with non functional. I chose this simplified way to deal with the imbalanced dataset not only due to time-constraints but because it is extremely difficult to get meaningful predictive power for the 'functional needs repair' label due to the nature of the dataset.

For the modeling in this notebook the target data will be relabeled as the following after binning function needs repair with non-functional:

non_functional: 0

functional: 1

In my preprocessing steps I had to use a work around for LabelEncoder to play nice. There is no value mapper method and it labels strings alphabetically. Therefore, I temporarily changed the status group labels to get encoded in the order I want them to be encoded.

```
In [4]: 1 df_binary['status_group'].value_counts(normalize=True)
```

```
Out[4]: Zfunctional      0.543081
Anon functional      0.456919
Name: status_group, dtype: float64
```

```
In [5]: 1 # le = LabelEncoder()
2
3 # le.fit(['functional', 'functional needs repair', 'non functional'])
4
5 # df['status_group'] = le.transform(df.status_group)
6
7 # df['status_group'].value_counts(normalize=True)
```



```
In [6]: 1 le = LabelEncoder()
        2
        3 le.fit(['Anon functional','Zfunctional'])
        4
        5 df_binary['status_group'] = le.transform(df_binary.status_group)
        6
        7 df_binary['status_group'].value_counts(normalize=True)
```

```
Out[6]: 1    0.543081
        0    0.456919
        Name: status_group, dtype: float64
```

Split Data into Test and Training Sets

```
In [7]: 1 from sklearn.model_selection import train_test_split
        2
        3 y = df_binary['status_group']
        4 X = df_binary.drop(columns=['status_group'], axis=1)
        5
        6 X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.3,
```

Run Master Cleaning Function

```
In [8]: 1 value_maps = get_col_val_mapping(df_binary)
        2 X_train = clean_dataframe(X_train, value_maps)
        3 X_test = clean_dataframe(X_test, value_maps)
        4 display(X_train.shape)
        5 X_test.shape
```

```
(41580, 239)
```

```
Out[8]: (17820, 239)
```

Data Modeling

Baseline model

I decided to use logistic regression as my baseline model. I chose to use an arbitrary large value for C and set the solver to 'liblinear.' I did not choose to fit an intercept for the baseline model.

```
In [9]: 1 # Instantiate the model
2 logreg = LogisticRegression(fit_intercept=False, C=1e12, solver='liblin
3
4 # Fit the model
5 logreg.fit(X_train, y_train)
6 y_hat_train = logreg.predict(X_train)
7 y_hat_test = logreg.predict(X_test)
8
9 print('Training Precision: ', precision_score(y_train, y_hat_train))
10 print('Testing Precision: ', precision_score(y_test, y_hat_test))
11 print('\n\n')
12
13 print('Training Recall: ', recall_score(y_train, y_hat_train))
14 print('Testing Recall: ', recall_score(y_test, y_hat_test))
15 print('\n\n')
16
17 print('Training Accuracy: ', accuracy_score(y_train, y_hat_train))
18 print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test))
19 print('\n\n')
20
21 print('Training F1-Score: ', f1_score(y_train, y_hat_train))
22 print('Testing F1-Score: ', f1_score(y_test, y_hat_test))
```

Training Precision: 0.7527743044701469
Testing Precision: 0.7528620752316918

Training Recall: 0.8548923896161527
Testing Recall: 0.8521184697655286

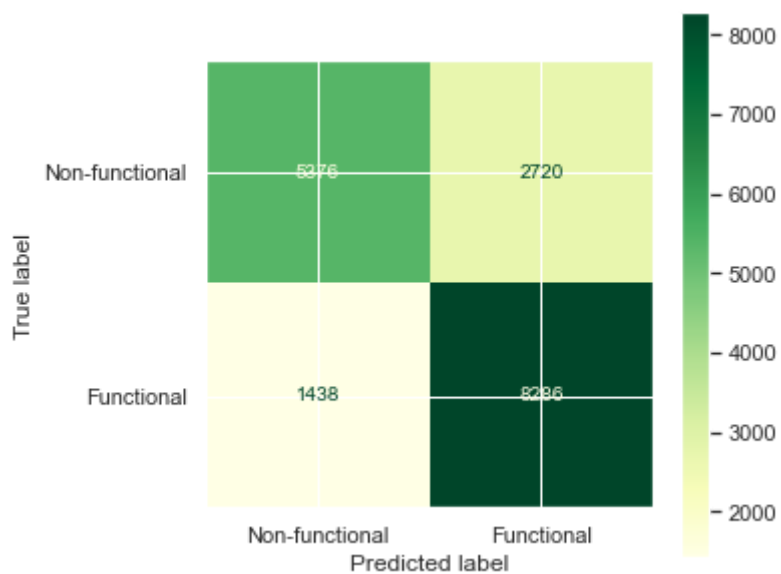
Training Accuracy: 0.7691919191919192
Testing Accuracy: 0.7666666666666667

Training F1-Score: 0.8005901053462713
Testing F1-Score: 0.7994211287988422


```
In [10]: 1 #Get the confusion matrix
2 lr_base_matrix = confusion_matrix(y_test, y_hat_test)
3 print(lr_base_matrix)
```

```
[[5376 2720]
 [1438 8286]]
```

```
In [11]: 1 fig_labels = ['Non-functional', 'Functional']
2 fig, ax = plt.subplots(figsize=(5, 5))
3
4 plot_confusion_matrix(logreg,
5                       X_test,
6                       y_test,
7                       ax=ax,
8                       cmap='YlGn',
9                       display_labels=fig_labels)
10
11 plt.show();
```



```
In [14]: 1 #confusion matrix percentages
2 fig, ax = plt.subplots(figsize=(5, 5))
3
4
5 sns.heatmap(lr_base_matrix/np.sum(lr_base_matrix), annot=True,
6             fmt='.2%', cmap='YlGn', ax=ax)
7
8 plt.show();
```



```
1 The baseline logisitc regression model had fairly decent scores all
2 around but I believe they can be improved.
3 Below a random grid search was used to create 300 of 600 possible
4 combinations to narrow down the paramters to do a full brute force
5 gridCV search. The score used for determining the 'best" model was
6 precision.
7 This is clearly not the way to build the model. Although I am greatly
8 concerned with increasing my precision in order to reduce false
9 positives, false negatives still have a fairly high cost as well.
10
11 The precision of the model is excellent. However, the accuracy and F1
12 scores are unacceptable.
13
14 I will describe why this is a poor model in more detail below under
15 the display of the confusion matrix.
```

```
In [15]: 1 start = time.time()
2
3 classifier_penalties = ['l1', 'l2']
4 classifier_Cs = np.logspace(-5, 5, 100)
5 classifier__solver = ['liblinear']
6
7 random_grid = {'penalty' : classifier_penalties,
8 'C' : classifier_Cs,
9 'solver' : classifier__solver}
10
11 lr = LogisticRegression()
12 # Random search of parameters, using 3 fold cross validation,
13 # search across 100 different combinations, and use all available cores
14 lr_random_precision = RandomizedSearchCV(scoring='precision',
15                                           estimator = lr,
16                                           param_distributions = random_grid,
17                                           n_iter = 100, cv = 3, verbose=2,
18                                           random_state=42,
19                                           n_jobs = -1)
20 # Fit the random search model
21 lr_random_precision.fit(X_train, y_train)
22
23
24
25 end = time.time()
26 print(end - start)
27
28 lr_best_rand_precision = lr_random_precision.best_estimator_
29 lr_best_rand_precision.fit(X_train,y_train)
30 y_hat_train_rp = lr_best_rand_precision.predict(X_train)
31 y_hat_test_rp = lr_best_rand_precision.predict(X_test)
```

Fitting 3 folds for each of 100 candidates, totalling 300 fits
245.77547788619995

```
In [16]: 1 print('Training Precision: ', precision_score(y_train, y_hat_train_rp))
2 print('Testing Precision: ', precision_score(y_test, y_hat_test_rp))
3 print('\n\n')
4
5
6 print('Training Accuracy: ', accuracy_score(y_train, y_hat_train_rp))
7 print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test_rp))
8 print('\n\n')
9
10 print('Training F1-Score: ', f1_score(y_train, y_hat_train_rp))
11 print('Testing F1-Score: ', f1_score(y_test, y_hat_test_rp))
```

Training Precision: 0.7525507212384192

Testing Precision: 0.7529572338489536

Training Accuracy: 0.7687830687830688

Testing Accuracy: 0.7663299663299663


Training F1-Score: 0.8001912045889101

Testing F1-Score: 0.7989765376074153

```
In [17]: 1 #Get the confusion matrix
2 lr_rand_prec_matrix = confusion_matrix(y_test, y_hat_test_rp)
3 print(lr_rand_prec_matrix)
```

```
[[5381 2715]
 [1449 8275]]
```

```
In [18]: 1 # Visualize your confusion matrix
2 fig, ax = plt.subplots(figsize=(5,5))
3
4
5 sns.heatmap(lr_rand_prec_matrix/np.sum(lr_rand_prec_matrix), annot=True,
6             fmt='.2%', cmap='YlGn', ax=ax)
7
8 plt.show();
```



The model did not perform much different than before the random search. Perhaps scoring on precision is not the way to go, as we also want to lower false negatives.

I will try scoring on `f1_score` as my target for iterative modeling as well as precision and compare.

Below I decided to run a full grid search with the same parameters as my random grid, only this time with focus on improving `f1_score`.


```

In [19]: 1 start = time.time()
2
3 classifier_penalties = ['l1', 'l2']
4 classifier_Cs = np.logspace(-5, 5, 100)
5 classifier__solver = ['liblinear']
6
7 lr_grid = {'penalty' : classifier_penalties,
8           'C' : classifier_Cs,
9           'solver' : classifier__solver
10          }
11
12 lr = LogisticRegression()
13 # Random search of parameters, using 3 fold cross validation,
14 # search across 100 different combinations, and use all available cores
15 lr_grid_f1 = GridSearchCV(estimator=lr,
16                           param_grid=lr_grid,
17                           scoring='f1',
18                           cv=3,
19                           n_jobs = -1,
20                           verbose=2
21                          )
22
23 # RandomizedSearchCV(estimator = lr,
24 #                    param_distributions = random_grid,
25 #                    n_iter = 100, cv = 3, verbose=2,
26 #                    random_state=42,
27 #                    n_jobs = -1)
28 # Fit the random search model
29 lr_grid_f1.fit(X_train, y_train)
30
31
32 end = time.time()
33 print(end - start)
34
35
36 # lr_best_grid_f1 = lr_grid_f1.best_estimator_
37 # lr_best_grid_f1.fit(X_train,y_train)
38 # y_hat_train_gs = lr_best_grid_f1.predict(X_train)
39 # y_hat_test_gs = lr_best_grid_f1.predict(X_test)

```

Fitting 3 folds for each of 200 candidates, totalling 600 fits
447.88261699676514

```

In [20]: 1 lr_grid_f1.best_estimator_

```

```

Out[20]: LogisticRegression(C=0.44306214575838776, penalty='l1', solver='liblinear')

```

```
In [21]: 1 lr_best_grid_f1 = lr_grid_f1.best_estimator_  
2 lr_best_grid_f1.fit(X_train,y_train)  
3 y_hat_train_gs = lr_best_grid_f1.predict(X_train)  
4 y_hat_test_gs = lr_best_grid_f1.predict(X_test)  
5  
6  
7 print('Training Precision: ', precision_score(y_train, y_hat_train_gs))  
8 print('Testing Precision: ', precision_score(y_test, y_hat_test_gs))  
9 print('\n\n')  
10  
11  
12 print('Training Accuracy: ', accuracy_score(y_train, y_hat_train_gs))  
13 print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test_gs))  
14 print('\n\n')  
15  
16 print('Training F1-Score: ', f1_score(y_train, y_hat_train_gs))  
17 print('Testing F1-Score: ', f1_score(y_test, y_hat_test_gs))
```


Training Precision: 0.7507864383082838
Testing Precision: 0.7509286943915919

Training Accuracy: 0.7686387686387687
Testing Accuracy: 0.7651515151515151


Training F1-Score: 0.8007621572363517
Testing F1-Score: 0.7984201146380233

```
In [22]: 1 #best grid search confusion matrix
2 lr_grid_f1_matrix = confusion_matrix(y_test, y_hat_test_gs)
3 print(lr_grid_f1_matrix)
4
5 # Visualize your confusion matrix
6 fig, ax = plt.subplots(figsize=(5,5))
7
8
9 sns.heatmap(lr_grid_f1_matrix/np.sum(lr_grid_f1_matrix), annot=True,
10             fmt='.2%', cmap='YlGn', ax=ax)
11
12 plt.show();
```


```
[[5347 2749]
 [1436 8288]]
```



```
In [23]: 1 #baseline confusion matrix
2 fig, ax = plt.subplots(figsize=(5,5))
3
4
5 sns.heatmap(lr_base_matrix/np.sum(lr_base_matrix), annot=True,
6             fmt='.2%', cmap='YlGn', ax=ax)
7
8 plt.show();
```



```
In [24]: 1 fig, ax = plt.subplots(figsize=(5,5))
2
3 plot_confusion_matrix(lr_best_grid_f1,
4                       X_test,
5                       y_test,
6                       ax=ax,
7                       cmap='YlGn',
8                       display_labels=fig_labels)
9
10 plt.show();
```



The gridsearch did not do much to improve the model.

Random Forest Classifier

Below I have created a grid of paramters in order to tune a random forest classifier. This grid will be used to perform a random grid search in order to narrow down the paramters to run a more robust grid search. The run time for a Random Forest has potential to be much higher than a logistic regression model. Therefore, a random search to narrow down the paramters for a grid search was carried out in order to reduce total run time while still keeping a fairly broad search.

```
In [25]: 1 # Number of trees in random forest
2 n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000, n
3 # Number of features to consider at every split
4 max_features = ['auto', 'sqrt']
5 # Maximum number of levels in tree
6 max_depth = [int(x) for x in np.linspace(10, 110, num = 11)]
7 max_depth.append(None)
8 # Minimum number of samples required to split a node
9 min_samples_split = [2, 5, 10]
10 # Minimum number of samples required at each leaf node
11 min_samples_leaf = [1, 2, 4]
12 # Create the random grid
13 random_grid = {'n_estimators': n_estimators,
14                 'max_features': max_features,
15                 'max_depth': max_depth,
16                 'min_samples_split': min_samples_split,
17                 'min_samples_leaf': min_samples_leaf}
18 random_grid
```

```
Out[25]: {'n_estimators': [200, 400, 600, 800, 1000, 1200, 1400, 1600, 1800, 2000],
'max_features': ['auto', 'sqrt'],
'max_depth': [10, 20, 30, 40, 50, 60, 70, 80, 90, 100, 110, None],
'min_samples_split': [2, 5, 10],
'min_samples_leaf': [1, 2, 4]}
```

```

In [26]: 1 start = time.time()
2
3 # Use the random grid to search for best hyperparameters
4 # First create the base model to tune
5 rfc = RandomForestClassifier()
6 # Random search of parameters, using 3 fold cross validation,
7 # search across 100 different combinations, and use all available cores
8 rfc_random = RandomizedSearchCV(scoring='f1',
9                                 estimator = rfc,
10                                param_distributions = random_grid,
11                                n_iter = 33,
12                                cv = 3,
13                                verbose=2,
14                                random_state=42,
15                                n_jobs = -1)
16 # Fit the random search model
17
18 rfc_random.fit(X_train,y_train)
19 rfc_best_random = rfc_random.best_estimator_
20 rfc_best_random.fit(X_train,y_train)
21 y_hat_train_rrf = rfc_best_random.predict(X_train)
22 y_hat_test_rrf = rfc_best_random.predict(X_test)
23
24
25 end = time.time()
26 print(end - start)
27
28 print('Training Precision: ', precision_score(y_train, y_hat_train_rrf))
29 print('Testing Precision: ', precision_score(y_test, y_hat_test_rrf))
30 print('\n\n')
31
32
33 print('Training Accuracy: ', accuracy_score(y_train, y_hat_train_rrf))
34 print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test_rrf))
35 print('\n\n')
36
37 print('Training F1-Score: ', f1_score(y_train, y_hat_train_rrf))
38 print('Testing F1-Score: ', f1_score(y_test, y_hat_test_rrf))

```

Fitting 3 folds for each of 33 candidates, totalling 99 fits

1702.4333066940308

Training Precision: 0.8238002619775334

Testing Precision: 0.7885450346420323

Training Accuracy: 0.8504088504088504

Testing Accuracy: 0.8048821548821549

Training F1-Score: 0.8696781763325511


Testing F1-Score: 0.8307946858727917

```
In [27]: 1 rfc_random.best_estimator_.get_params()
```

```
Out[27]: {'bootstrap': True,
'ccp_alpha': 0.0,
'class_weight': None,
'criterion': 'gini',
'max_depth': 40,
'max_features': 'sqrt',
'max_leaf_nodes': None,
'max_samples': None,
'min_impurity_decrease': 0.0,
'min_impurity_split': None,
'min_samples_leaf': 2,
'min_samples_split': 5,
'min_weight_fraction_leaf': 0.0,
'n_estimators': 1000,
'n_jobs': None,
'oob_score': False,
'random_state': None,
'verbose': 0,
'warm_start': False}
```

```
In [28]: 1 #best random search confusion matrix for Random Forest
2 lr_grid_f1_matrix_rrf = confusion_matrix(y_test, y_hat_test_rrf)
3 print(lr_grid_f1_matrix_rrf)
4
5 # Visualize your confusion matrix
6 fig, ax = plt.subplots(figsize=(5,5))
7
8
9 sns.heatmap(lr_grid_f1_matrix_rrf/np.sum(lr_grid_f1_matrix_rrf), annot=
10             fmt='.2%', cmap='YlGn', ax=ax)
11
12 plt.show();
```

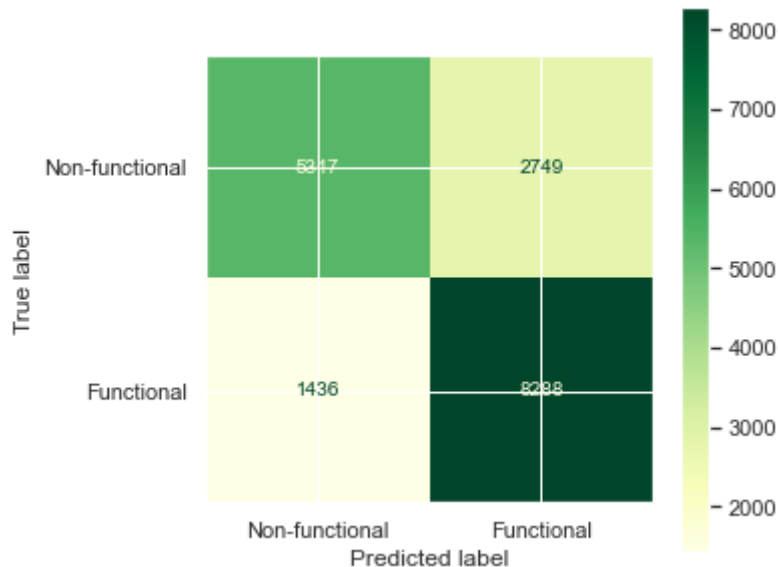
```
[[5807 2289]
 [1188 8536]]
```




```

In [29]: 1 fig, ax = plt.subplots(figsize=(5,5))
          2
          3 plot_confusion_matrix(lr_best_grid_f1,
          4                             X_test,
          5                             y_test,
          6                             ax=ax,
          7                             cmap='YlGn',
          8                             display_labels=fig_labels)
          9
          10 plt.show();

```



This random forest classifier outperforms the best logistic regression model in every relevant metric. It improves upon the model accuracy and precision while also lowering the occurrence false negatives.

A grid search will be performed below to see if any improvements can be made to model performance.

```

In [30]: 1 # Number of trees in random forest
          2 n_estimators = [500, 600, 700]
          3 # Number of features to consider at every split
          4 max_features = [14,15,16]
          5 # Maximum number of levels in tree
          6 # Minimum number of samples required to split a node
          7 min_samples_split = [8,10,12]
          8 # Minimum number of samples required at each leaf node
          9 # Create the random grid
          10 cv_grid = {'n_estimators': n_estimators,
          11                'max_features': max_features,
          12                'min_samples_split': min_samples_split}
          13
          14 cv_grid

```

```

Out[30]: {'n_estimators': [500, 600, 700],
          'max_features': [14, 15, 16],
          'min_samples_split': [8, 10, 12]}

```

```

In [31]: 1 start = time.time()
2
3 rfc = RandomForestClassifier()
4 # Instantiate the grid search model
5 rfc_gs = GridSearchCV(scoring='f1',
6                       estimator = rfc,
7                       param_grid = cv_grid,
8                       cv = 3,
9                       n_jobs = -1,
10                      verbose = 2)
11
12
13 rfc_gs.fit(X_train,y_train)
14 rfc_best_gs = rfc_gs.best_estimator_
15 rfc_best_gs.fit(X_train,y_train)
16 y_hat_train_grf = rfc_best_gs.predict(X_train)
17 y_hat_test_grf = rfc_best_gs.predict(X_test)
18
19
20 end = time.time()
21 print(end - start)
22
23 print('Training Precision: ', precision_score(y_train, y_hat_train_grf))
24 print('Testing Precision: ', precision_score(y_test, y_hat_test_grf))
25 print('\n\n')
26
27
28 print('Training Accuracy: ', accuracy_score(y_train, y_hat_train_grf))
29 print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test_grf))
30 print('\n\n')
31
32 print('Training F1-Score: ', f1_score(y_train, y_hat_train_grf))
33 print('Testing F1-Score: ', f1_score(y_test, y_hat_test_grf))

```

Fitting 3 folds for each of 27 candidates, totalling 81 fits

958.5107200145721

Training Precision: 0.8589955394019494

Testing Precision: 0.7988826815642458

Training Accuracy: 0.8761183261183261

Testing Accuracy: 0.8025813692480359

Training F1-Score: 0.889811110873425

Testing F1-Score: 0.8250273550184024

Random forest models using random and grid search CV with scoring set to 'f1'

Testing random search Precision: 0.7885450346420323

Testing random search Accuracy: 0.8048821548821549

Testing random search F1-Score: 0.8307946858727917

[5807 2289]

[1188 8536]

vs

Testing grid search Precision: 0.7988826815642458

Testing grid search Accuracy: 0.8025813692480359

Testing grid search F1-Score: 0.8250273550184024

[6008 2088]

[1430 8294]

The grid-search-tuned random forest classifier correctly labeled about 200 more non-functional pumps and had about 200 less false positives than the random-search-tuned random forest classifier.

I believe the higher precision of the grid search random forest classifier is more important than the slight hit to the f1-score and accuracy.

```


In [66]: 1 #best grid search confusion matrix for Random Forest
2 lr_grid_fl_matrix_grf = confusion_matrix(y_test, y_hat_test_grf)
3 print(lr_grid_fl_matrix_grf)
4
5 # Visualize your confusion matrix
6 fig, ax = plt.subplots(figsize=(5, 5))
7
8
9 sns.heatmap(lr_grid_fl_matrix_grf/np.sum(lr_grid_fl_matrix_grf),
10             annot=True,
11             fmt='.2%',
12             cmap='YlGn',
13             ax=ax)
14
15 plt.show();
16
17 #current best model

```

```

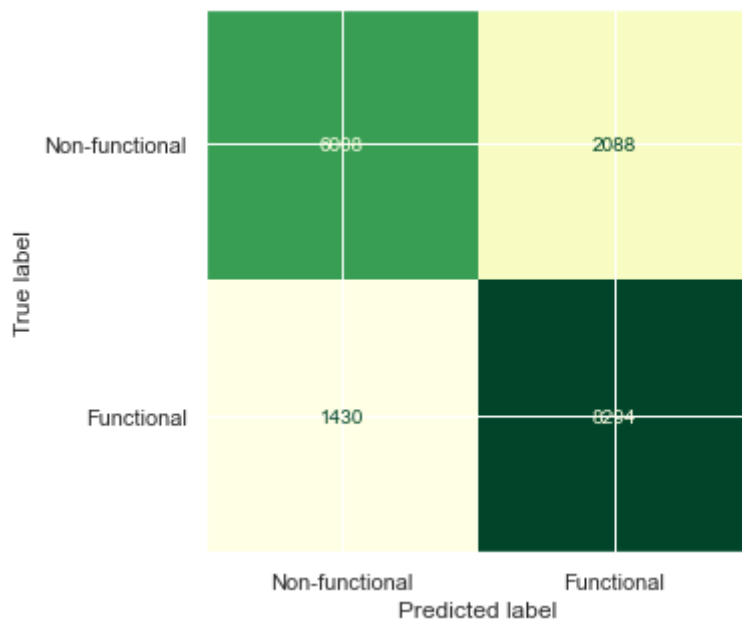
[[ 6008  2088]
 [ 1430  8294]]

```



The results of the grid search reduced the number of false positives, which is great, but it slightly increased the number of false negatives. It identified a similar number of functional pumps as the other random forest model, but it also identified slightly more non-functional pumps correctly. Performance was slightly improved overall.

```
In [68]: 1 #current best model
2
3 fig, ax = plt.subplots(figsize=(5, 5))
4
5 plot_confusion_matrix(rfc_best_gs,
6                       X_test,
7                       y_test,
8                       ax=ax,
9                       cmap='YlGn',
10                      colorbar=False,
11                      display_labels=fig_labels)
12
13 plt.show();
```



```

In [57]: 1 start = time.time()
2
3 # Use the random grid to search for best hyperparameters
4 # First create the base model to tune
5 rfc = RandomForestClassifier()
6 # Random search of parameters, using 3 fold cross validation,
7 # search across 100 different combinations, and use all available cores
8 rfc_random2 = RandomizedSearchCV(scoring='precision',
9                                 estimator = rfc,
10                                param_distributions = random_grid,
11                                n_iter = 33,
12                                cv = 3,
13                                verbose=2,
14                                random_state=42,
15                                n_jobs = -1)
16 # Fit the random search model
17
18 rfc_random2.fit(X_train,y_train)
19 rfc_best_random2 = rfc_random2.best_estimator_
20 rfc_best_random2.fit(X_train,y_train)
21 y_hat_train_rrf2 = rfc_best_random2.predict(X_train)
22 y_hat_test_rrf2 = rfc_best_random2.predict(X_test)
23
24
25 end = time.time()
26 print(end - start)
27
28 print('Training Precision: ', precision_score(y_train, y_hat_train_rrf2))
29 print('Testing Precision: ', precision_score(y_test, y_hat_test_rrf2))
30 print('\n\n')
31
32
33 print('Training Accuracy: ', accuracy_score(y_train, y_hat_train_rrf2))
34 print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test_rrf2))
35 print('\n\n')
36
37 print('Training F1-Score: ', f1_score(y_train, y_hat_train_rrf2))
38 print('Testing F1-Score: ', f1_score(y_test, y_hat_test_rrf2))

```

Fitting 3 folds for each of 33 candidates, totalling 99 fits

1520.5148952007294

Training Precision: 0.878849845588849

Testing Precision: 0.8001746047143273

Training Accuracy: 0.8946849446849446

Testing Accuracy: 0.8016273849607183

Training F1-Score: 0.9058218809815688


Testing F1-Score: 0.8235411570908002

In [54]: 1 rfc_best_gs2

Out[54]: RandomForestClassifier(max_features=16, min_samples_split=8, n_estimators=500)

```
In [70]: 1 #best grid search confusion matrix for Random Forest
2 lr_grid_p_matrix_grf = confusion_matrix(y_test, y_hat_test_rrf2)
3 print(lr_grid_p_matrix_grf)
4
5 # Visualize your confusion matrix
6 fig, ax = plt.subplots(figsize=(5, 5))
7
8
9 sns.heatmap(lr_grid_p_matrix_grf/np.sum(lr_grid_p_matrix_grf),
10             annot=True,
11             fmt='.2%',
12             cmap='YlGn',
13             ax=ax)
14
15 plt.show();
```

```
[[6036 2060]
 [1475 8249]]
```



current model (random search precision scoring)

Testing Precision: 0.8001746047143273

Testing Accuracy: 0.8016273849607183

Testing F1-Score: 0.8235411570908002

[6036, 2060]

[1475, 8249]

vs

best model so far (grid search f1 scoring)

Testing grid search Precision: 0.7988826815642458

Testing grid search Accuracy: 0.8025813692480359

Testing grid search F1-Score: 0.8250273550184024

[6008, 2088]

[1430, 8294]

Both models perform very similarly. The current model has slightly less false positives than the best model so far.

```
In [71]: 1 # Number of trees in random forest
2 n_estimators = [500, 600, 700]
3 # Number of features to consider at every split
4 max_features = [16,17,18]
5 # Maximum number of levels in tree
6 # Minimum number of samples required to split a node
7 min_samples_split = [8,10,12]
8 # Minimum number of samples required at each leaf node
9 # Create the random grid
10 cv_grid2 = {'n_estimators': n_estimators,
11             'max_features': max_features,
12             'min_samples_split': min_samples_split}
13
14 cv_grid2
```

```
Out[71]: {'n_estimators': [500, 600, 700],
'max_features': [16, 17, 18],
'min_samples_split': [8, 10, 12]}
```

```

In [72]: 1 start = time.time()
2
3 rfc = RandomForestClassifier()
4 # Instantiate the grid search model
5 rfc_gs2 = GridSearchCV(scoring='precision',
6                         estimator = rfc,
7                         param_grid = cv_grid2,
8                         cv = 3,
9                         n_jobs = -1,
10                        verbose = 2)
11
12
13 rfc_gs2.fit(X_train,y_train)
14 rfc_best_gs2 = rfc_gs2.best_estimator_
15 rfc_best_gs2.fit(X_train,y_train)
16 y_hat_train_grf2 = rfc_best_gs2.predict(X_train)
17 y_hat_test_grf2 = rfc_best_gs2.predict(X_test)
18
19
20 end = time.time()
21 print(end - start)
22
23 print('Training Precision: ', precision_score(y_train, y_hat_train_grf2))
24 print('Testing Precision: ', precision_score(y_test, y_hat_test_grf2))
25 print('\n\n')
26
27
28 print('Training Accuracy: ', accuracy_score(y_train, y_hat_train_grf2))
29 print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test_grf2))
30 print('\n\n')
31
32 print('Training F1-Score: ', f1_score(y_train, y_hat_train_grf2))
33 print('Testing F1-Score: ', f1_score(y_test, y_hat_test_grf2))

```

Fitting 3 folds for each of 27 candidates, totalling 81 fits

1050.9529011249542

Training Precision: 0.8690639127369251

Testing Precision: 0.8000386025863733

Training Accuracy: 0.8860509860509861

Testing Accuracy: 0.8032547699214366

Training F1-Score: 0.8984264459975132

Testing F1-Score: 0.825450562580902

```


In [74]: 1 #best grid search confusion matrix for Random Forest
2 lr_grid_p_matrix_grf2 = confusion_matrix(y_test, y_hat_test_grf2)
3 print(lr_grid_p_matrix_grf2)
4
5 # Visualize your confusion matrix
6 fig, ax = plt.subplots(figsize=(5, 5))
7
8
9 sns.heatmap(lr_grid_p_matrix_grf2/np.sum(lr_grid_p_matrix_grf2),
10             annot=True,
11             fmt='.2%',
12             cmap='YlGn',
13             ax=ax)
14
15 plt.show();

```

```

[[6024 2072]
 [1434 8290]]

```



I believe the random forest classifier using the hyperparameters determined through the grid search using precision as the tuning parameter produced the best model as a balance between accuracy and low false positives and false negatives.

grid search precision scoring (best model)

Testing Precision: 0.8000386025863733

Testing Accuracy: 0.8032547699214366

Testing F1-Score: 0.825450562580902 [6024, 2072]
[1434, 8290]

random search presicion scoring

Testing Precision: 0.8001746047143273

Testing Accuracy: 0.8016273849607183

Testing F1-Score: 0.8235411570908002

[6036, 2060]

[1475, 8249]

vs

grid search f1 scoring

Testing grid search Precision: 0.7988826815642458

Testing grid search Accuracy: 0.8025813692480359

Testing grid search F1-Score: 0.8250273550184024

[6008, 2088]

[1430, 8294]

XGBoost Classification Iterative Modeling

A similar method of iterative modeling was carried out using XGBoost as was used for tuning the random forest classifier: a random grid cv will be used to reduce the number of paramters to perform in a more robust grid search. A baseline XGBoost classifier was first created using default parameters before iterative modeling/hyper-paramter tuning was carried out.

```
In [35]: 1 # Instantiate the model
2 xgc_base = XGBClassifier()
3
4 # Fit the model
5 xgc_base.fit(X_train, y_train)
6 y_hat_train = xgc_base.predict(X_train)
7 y_hat_test = xgc_base.predict(X_test)
8
9 print('Training Precision: ', precision_score(y_train, y_hat_train))
10 print('Testing Precision: ', precision_score(y_test, y_hat_test))
11 print('\n\n')
12
13 print('Training Recall: ', recall_score(y_train, y_hat_train))
14 print('Testing Recall: ', recall_score(y_test, y_hat_test))
15 print('\n\n')
16
17 print('Training Accuracy: ', accuracy_score(y_train, y_hat_train))
18 print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test))
19 print('\n\n')
20
21 print('Training F1-Score: ', f1_score(y_train, y_hat_train))
22 print('Testing F1-Score: ', f1_score(y_test, y_hat_test))
```

Training Precision: 0.7247322505142209
Testing Precision: 0.7266787958981145


Training Recall: 0.9068560017750167
Testing Recall: 0.9036404771698889

Training Accuracy: 0.7628427128427129
Testing Accuracy: 0.7619528619528619

Training F1-Score: 0.8056294719413399
Testing F1-Score: 0.8055555555555556

```
In [36]: 1 #baseline XGBoost model
2 xgs_base_matrix = confusion_matrix(y_test, y_hat_test)
3 print(xgs_base_matrix)
4
5 # Visualize your confusion matrix
6 fig, ax = plt.subplots(figsize=(5, 5))
7
8
9 sns.heatmap(xgs_base_matrix/np.sum(xgs_base_matrix), annot=True,
10             fmt='.2%', cmap='YlGn', ax=ax)
11
12 plt.show();
```

```
[[4791 3305]
 [ 937 8787]]
```



The baseline XGBoost model has a low number of false negatives, but a relatively high number of false positives. It also has a lower accuracy than either random forest model. Some hyper-parameter tuning will be carried out below in an attempt to improve the model.

```

In [37]: 1 params = {
2         'min_child_weight': [1, 5, 10],
3         'gamma': [0.5, 1, 1.5, 2, 5],
4         'subsample': [0.6, 0.8, 1.0],
5         'colsample_bytree': [0.6, 0.8, 1.0],
6         'max_depth': [3, 4, 5]
7     }
8
9     xgc = XGBClassifier()
10    # Random search of parameters, using 3 fold cross validation,
11    # search across 100 different combinations, and use all available cores
12    xgc_random = RandomizedSearchCV(scoring='f1',
13                                    estimator = xgc,
14                                    param_distributions = params,
15                                    n_iter = 33, cv = 3, verbose=2,
16                                    random_state=42,
17                                    n_jobs = -1)
18    # Fit the random search model
19    xgc_random.fit(X_train, y_train)
20
21
22    end = time.time()
23    print(end - start)
24
25
26    # xgc_best_random = xgc_random.best_estimator_
27    # xgc_best_random.fit(X_train,y_train)
28    # y_hat_train_brm = xgc_best_random.predict(X_train)
29    # y_hat_test_brm = xgc_best_random.predict(X_test)
30
31    # print('Training Precision: ', precision_score(y_train, y_hat_train_brm))
32    # print('Testing Precision: ', precision_score(y_test, y_hat_test_brm))
33    # print('\n\n')
34
35    # print('Training Recall: ', recall_score(y_train, y_hat_train_brm))
36    # print('Testing Recall: ', recall_score(y_test, y_hat_test_brm))
37    # print('\n\n')
38
39    # print('Training Accuracy: ', accuracy_score(y_train, y_hat_train_brm))
40    # print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test_brm))
41    # print('\n\n')
42
43    # print('Training F1-Score: ', f1_score(y_train, y_hat_train_brm))
44    # print('Testing F1-Score: ', f1_score(y_test, y_hat_test_brm))

```

Fitting 3 folds for each of 33 candidates, totalling 99 fits
 1733.5508239269257


```
In [38]: 1 xgc_best_random = xgc_random.best_estimator_  
2 xgc_best_random.fit(X_train,y_train)  
3 y_hat_train_brm = xgc_best_random.predict(X_train)  
4 y_hat_test_brm = xgc_best_random.predict(X_test)  
5  
6 print('Training Precision: ', precision_score(y_train, y_hat_train_brm))  
7 print('Testing Precision: ', precision_score(y_test, y_hat_test_brm))  
8 print('\n\n')  
9  
10 print('Training Recall: ', recall_score(y_train, y_hat_train_brm))  
11 print('Testing Recall: ', recall_score(y_test, y_hat_test_brm))  
12 print('\n\n')  
13  
14 print('Training Accuracy: ', accuracy_score(y_train, y_hat_train_brm))  
15 print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test_brm))  
16 print('\n\n')  
17  
18 print('Training F1-Score: ', f1_score(y_train, y_hat_train_brm))  
19 print('Testing F1-Score: ', f1_score(y_test, y_hat_test_brm))
```

Training Precision: 0.7548979817732354
Testing Precision: 0.75055966936456

Training Recall: 0.9079210117594853
Testing Recall: 0.896441793500617

Training Accuracy: 0.7903318903318903
Testing Accuracy: 0.7809203142536476

Training F1-Score: 0.8243684274144808
Testing F1-Score: 0.8170400224950791

```


In [39]: 1 #the model performed better than the base XGBoost model
          2 #however it still underperforms compared to my best
          3 #random forest classifier
          4
          5 xgs_rand_matrix = confusion_matrix(y_test, y_hat_test_brm)
          6 print(xgs_rand_matrix)
          7
          8 # Visualize your confusion matrix
          9 fig, ax = plt.subplots(figsize=(5, 5))
         10
         11
         12 sns.heatmap(xgs_rand_matrix/np.sum(xgs_rand_matrix), annot=True,
         13           fmt='.2%', cmap='YlGn', ax=ax)
         14
         15 plt.show();

```

```

[[5199 2897]
 [1007 8717]]

```



Improvements to accuracy and f1-score. Still needs improvement to both to compete with the random forest models.

```

In [40]: 1 xgc_best_random

```

```

Out[40]: XGBClassifier(colsample_bytree=0.8, gamma=0.5, max_depth=5, subsample=1.0)

```

```

In [41]: 1 params2 = {
2         'gamma': [0.1, 0.3, 0.5],
3         'colsample_bytree': [0.8, 0.9],
4         'max_depth': [4, 5, 6]
5         }
6
7 start = time.time()
8
9 xgc2 = XGBClassifier()
10 # Random search of parameters, using 3 fold cross validation,
11 # search across 100 different combinations, and use all available cores
12 xgc_gs = GridSearchCV(scoring='f1',
13                       estimator = xgc2,
14                       param_grid = params2,
15                       cv = 3,
16                       verbose=2,
17                       n_jobs = -1)
18
19
20 xgc_gs.fit(X_train, y_train)
21 xgc_best_gs = xgc_gs.best_estimator_
22 xgc_best_gs.fit(X_train, y_train)
23 y_hat_train_bgs = xgc_best_gs.predict(X_train)
24 y_hat_test_bgs = xgc_best_gs.predict(X_test)
25
26
27 end = time.time()
28 print(end - start)
29
30 print('Training Precision: ', precision_score(y_train, y_hat_train_bgs))
31 print('Testing Precision: ', precision_score(y_test, y_hat_test_bgs))
32 print('\n\n')
33
34
35 print('Training Accuracy: ', accuracy_score(y_train, y_hat_train_bgs))
36 print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test_bgs))
37 print('\n\n')
38
39 print('Training F1-Score: ', f1_score(y_train, y_hat_train_bgs))
40 print('Testing F1-Score: ', f1_score(y_test, y_hat_test_bgs))

```

Fitting 3 folds for each of 18 candidates, totalling 54 fits

582.7580349445343

Training Precision: 0.7660139023843336

Testing Precision: 0.759965034965035

Training Accuracy: 0.8004088504088505

Testing Accuracy: 0.7881032547699215

Training F1-Score: 0.8316393808451504

Testing F1-Score: 0.8215838215838216

```


In [73]: 1 #slight improvement to f1-score as compared to before the grid-search
2 #still too many false positives to compete with random forest classifier
3 xgs_grid_matrix = confusion_matrix(y_test, y_hat_test_bgs)
4 print(xgs_grid_matrix)
5
6 # Visualize your confusion matrix
7 fig, ax = plt.subplots(figsize=(5, 5))
8
9
10 sns.heatmap(xgs_grid_matrix/np.sum(xgs_grid_matrix), annot=True,
11             fmt='.2%', cmap='YlGn', ax=ax)
12
13 plt.show();

```

```

[[5350 2746]
 [1030 8694]]

```



```

In [52]: 1 #change scoring to 'precision' for the random search
2 #will do the same for the grid search afterwards
3
4
5 params = {
6     'min_child_weight': [1, 5, 10],
7     'gamma': [0.5, 1, 1.5, 2, 5],
8     'subsample': [0.6, 0.8, 1.0],
9     'colsample_bytree': [0.6, 0.8, 1.0],
10    'max_depth': [3, 4, 5]
11    }
12
13 xgc = XGBClassifier()
14 # Random search of parameters, using 3 fold cross validation,
15 # search across 100 different combinations, and use all available cores
16 xgc_random2 = RandomizedSearchCV(scoring='precision',
17                                  estimator = xgc,
18                                  param_distributions = params,
19                                  n_iter = 33, cv = 3, verbose=2,
20                                  random_state=42,
21                                  n_jobs = -1)
22 # Fit the random search model
23 xgc_random2.fit(X_train, y_train)
24
25
26 end = time.time()
27 print(end - start)
28
29
30 xgc_best_random2 = xgc_random2.best_estimator_
31 xgc_best_random2.fit(X_train,y_train)
32 y_hat_train_brm2 = xgc_best_random2.predict(X_train)
33 y_hat_test_brm2 = xgc_best_random2.predict(X_test)
34
35 print('Training Precision: ', precision_score(y_train, y_hat_train_brm2))
36 print('Testing Precision: ', precision_score(y_test, y_hat_test_brm2))
37 print('\n\n')
38
39 print('Training Recall: ', recall_score(y_train, y_hat_train_brm2))
40 print('Testing Recall: ', recall_score(y_test, y_hat_test_brm2))
41 print('\n\n')
42
43 print('Training Accuracy: ', accuracy_score(y_train, y_hat_train_brm2))
44 print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test_brm2))
45 print('\n\n')
46
47 print('Training F1-Score: ', f1_score(y_train, y_hat_train_brm2))
48 print('Testing F1-Score: ', f1_score(y_test, y_hat_test_brm2))

```

Fitting 3 folds for each of 33 candidates, totalling 99 fits

1486.0256688594818

Training Precision: 0.7548979817732354

Testing Precision: 0.75055966936456

Training Recall: 0.9079210117594853

Testing Recall: 0.896441793500617

Training Accuracy: 0.7903318903318903

Testing Accuracy: 0.7809203142536476

Training F1-Score: 0.8243684274144808


Testing F1-Score: 0.8170400224950791

In [55]: 1 xgc_best_random2

Out[55]: XGBClassifier(colsample_bytree=0.8, gamma=0.5, max_depth=5, subsample=1.0)

```
In [75]: 1 #still not performing as well as random forest classifier
2
3
4 xgs_grid_matrix_p = confusion_matrix(y_test, y_hat_test_brm2)
5 print(xgs_grid_matrix_p)
6
7 # Visualize your confusion matrix
8 fig, ax = plt.subplots(figsize=(5, 5))
9
10
11 sns.heatmap(xgs_grid_matrix_p/np.sum(xgs_grid_matrix_p), annot=True,
12             fmt='.2%', cmap='YlGn', ax=ax)
13
14 plt.show();
```

```
[[5199 2897]
 [1007 8717]]
```



```

In [59]: 1 params_p = {
2         'gamma': [0.3,0.4,0.5],
3         'colsample_bytree': [0.8],
4         'max_depth': [5,8,10]
5         }
6
7 start = time.time()
8
9 xgc2 = XGBClassifier()
10 # Random search of parameters, using 3 fold cross validation,
11 # search across 100 different combinations, and use all available cores
12 xgc_gs2 = GridSearchCV(scoring='precision',
13                        estimator = xgc2,
14                        param_grid = params_p,
15                        cv = 3,
16                        verbose=2,
17                        n_jobs = -1)
18
19
20 xgc_gs2.fit(X_train, y_train)
21 xgc_best_gs2 = xgc_gs2.best_estimator_
22 xgc_best_gs2.fit(X_train,y_train)
23 y_hat_train_bgs2 = xgc_best_gs2.predict(X_train)
24 y_hat_test_bgs2 = xgc_best_gs2.predict(X_test)
25
26
27 end = time.time()
28 print(end - start)
29
30 print('Training Precision: ', precision_score(y_train, y_hat_train_bgs2))
31 print('Testing Precision: ', precision_score(y_test, y_hat_test_bgs2))
32 print('\n\n')
33
34
35 print('Training Accuracy: ', accuracy_score(y_train, y_hat_train_bgs2))
36 print('Testing Accuracy: ', accuracy_score(y_test, y_hat_test_bgs2))
37 print('\n\n')
38
39 print('Training F1-Score: ', f1_score(y_train, y_hat_train_bgs2))
40 print('Testing F1-Score: ', f1_score(y_test, y_hat_test_bgs2))

```

Fitting 3 folds for each of 9 candidates, totalling 27 fits

503.37574887275696

Training Precision: 0.8044286156704862

Testing Precision: 0.781864299302473

Training Accuracy: 0.8362914862914863


Testing Accuracy: 0.8035353535353535

Training F1-Score: 0.8592635474600451

Testing F1-Score: 0.8313827481577806

```
In [76]: 1 #decent model but not the best model so far
2
3 xgs_grid_matrix_gp = confusion_matrix(y_test, y_hat_test_bgs2)
4 print(xgs_grid_matrix_gp)
5
6 # Visualize your confusion matrix
7 fig, ax = plt.subplots(figsize=(5, 5))
8
9
10 sns.heatmap(xgs_grid_matrix_gp/np.sum(xgs_grid_matrix_gp), annot=True,
11             fmt='.2%', cmap='YlGn', ax=ax)
12
13 plt.show();
```

```
[[5688 2408]
 [1093 8631]]
```



The grid search produced a pretty good model, but my random forest classifier still performed better.

Evaluation


My random forest models outperformed my best logistic regression and XGBoost models in regards to the metrics that are most important given the business problem at hand.

The best random forest model had a great balance between accuracy, precision, and f1-score.


11.63% of pumps would be misclassified as functional using my best model. This means that 11.63% of the pumps would go untreated if this classifier was deployed to conduct predictive maintenance. However, it correctly identifies a high number of functional pumps correctly, which would save a lot of valuable resources, time and money, and it also identifies a large number of non-functional pumps correctly. Only 8.05% of functional pumps would be incorrectly identified as non-functional. This is the resource/time/money sink of my model, so keeping it so low is great.


```
In [90]: 1 fig, ax = plt.subplots(figsize=(10, 10))
2 print(lr_grid_p_matrix_grf2)
3
4 sns.heatmap(lr_grid_p_matrix_grf2/np.sum(lr_grid_p_matrix_grf2),
5             annot=True,
6             fmt='.2%',
7             cmap='YlGn',
8             cbar=False,
9             xticklabels=['Non-Functional', 'Functional'],
10            yticklabels=['Non-Functional', 'Functional'],
11            ax=ax)
12 plt.savefig('images/best_rfc_matrix')
13 plt.show();
```

```
[[6024 2072]
 [1434 8290]]
```




```
In [91]: 1 fig, ax = plt.subplots(figsize=(10, 10))
2
3 sns.heatmap(lr_grid_p_matrix_grf2,
4             annot=True,
5             fmt='n',
6             cmap='YlGn',
7             cbar=False,
8             xticklabels=['Non-Functional', 'Functional'],
9             yticklabels=['Non-Functional', 'Functional'],
10            ax=ax)
11 plt.savefig('images/best_rfc_matrix2')
12 plt.show();
```




Conclusions

I believe that my best classification model provides a powerful enough predictive ability to prove very valuable to the Ministry of Water. The amount of resources saved, the relatively low number of misclassified functional pumps, and the elimination of the need to physically sweep the functionality of all pumps can bring access to potable drinking water to a larger number of communities than before without predictive maintenance.

```
In [93]: 1 fig_matrix_60k = 60000*(lr_grid_p_matrix_grf2/np.sum(lr_grid_p_matrix_g
2 fig_matrix_10k = 10000*(lr_grid_p_matrix_grf2/np.sum(lr_grid_p_matrix_g
3
4 fig, ax = plt.subplots(figsize=(10, 10))
5
6
7 sns.heatmap(fig_matrix_60k,
8             annot=True,
9             fmt='.0f',
10            cmap='YlGn',
11            cbar=False,
12            ax=ax,
13            xticklabels=['Functional', 'Non-Functional'],
14            yticklabels=['Functional', 'Non-Functional']
15        )
16 plt.savefig('images/best_rfc_matrix_60k')
17 plt.show();
```



```
In [94]: 1 fig, ax = plt.subplots(figsize=(10, 10))
2
3
4 sns.heatmap(fig_matrix_10k,
5             annot=True,
6             fmt='.0f',
7             cmap='YlGn',
8             cbar=False,
9             ax=ax,
10            xticklabels=['Functional', 'Non-Functional'],
11            yticklabels=['Functional', 'Non-Functional']
12        )
13 plt.savefig('images/best_rfc_matrix_10k')
14 plt.show();
```



Thank you! For questions or comments please feel free to reach me by email.

Author: Dylan Dey

Email: ddey2985@gmail.com (<mailto:ddey2985@gmail.com>).

github: https://github.com/ddey117/Tanzanian_Water_Pump_Classification
(https://github.com/ddey117/Tanzanian_Water_Pump_Classification)